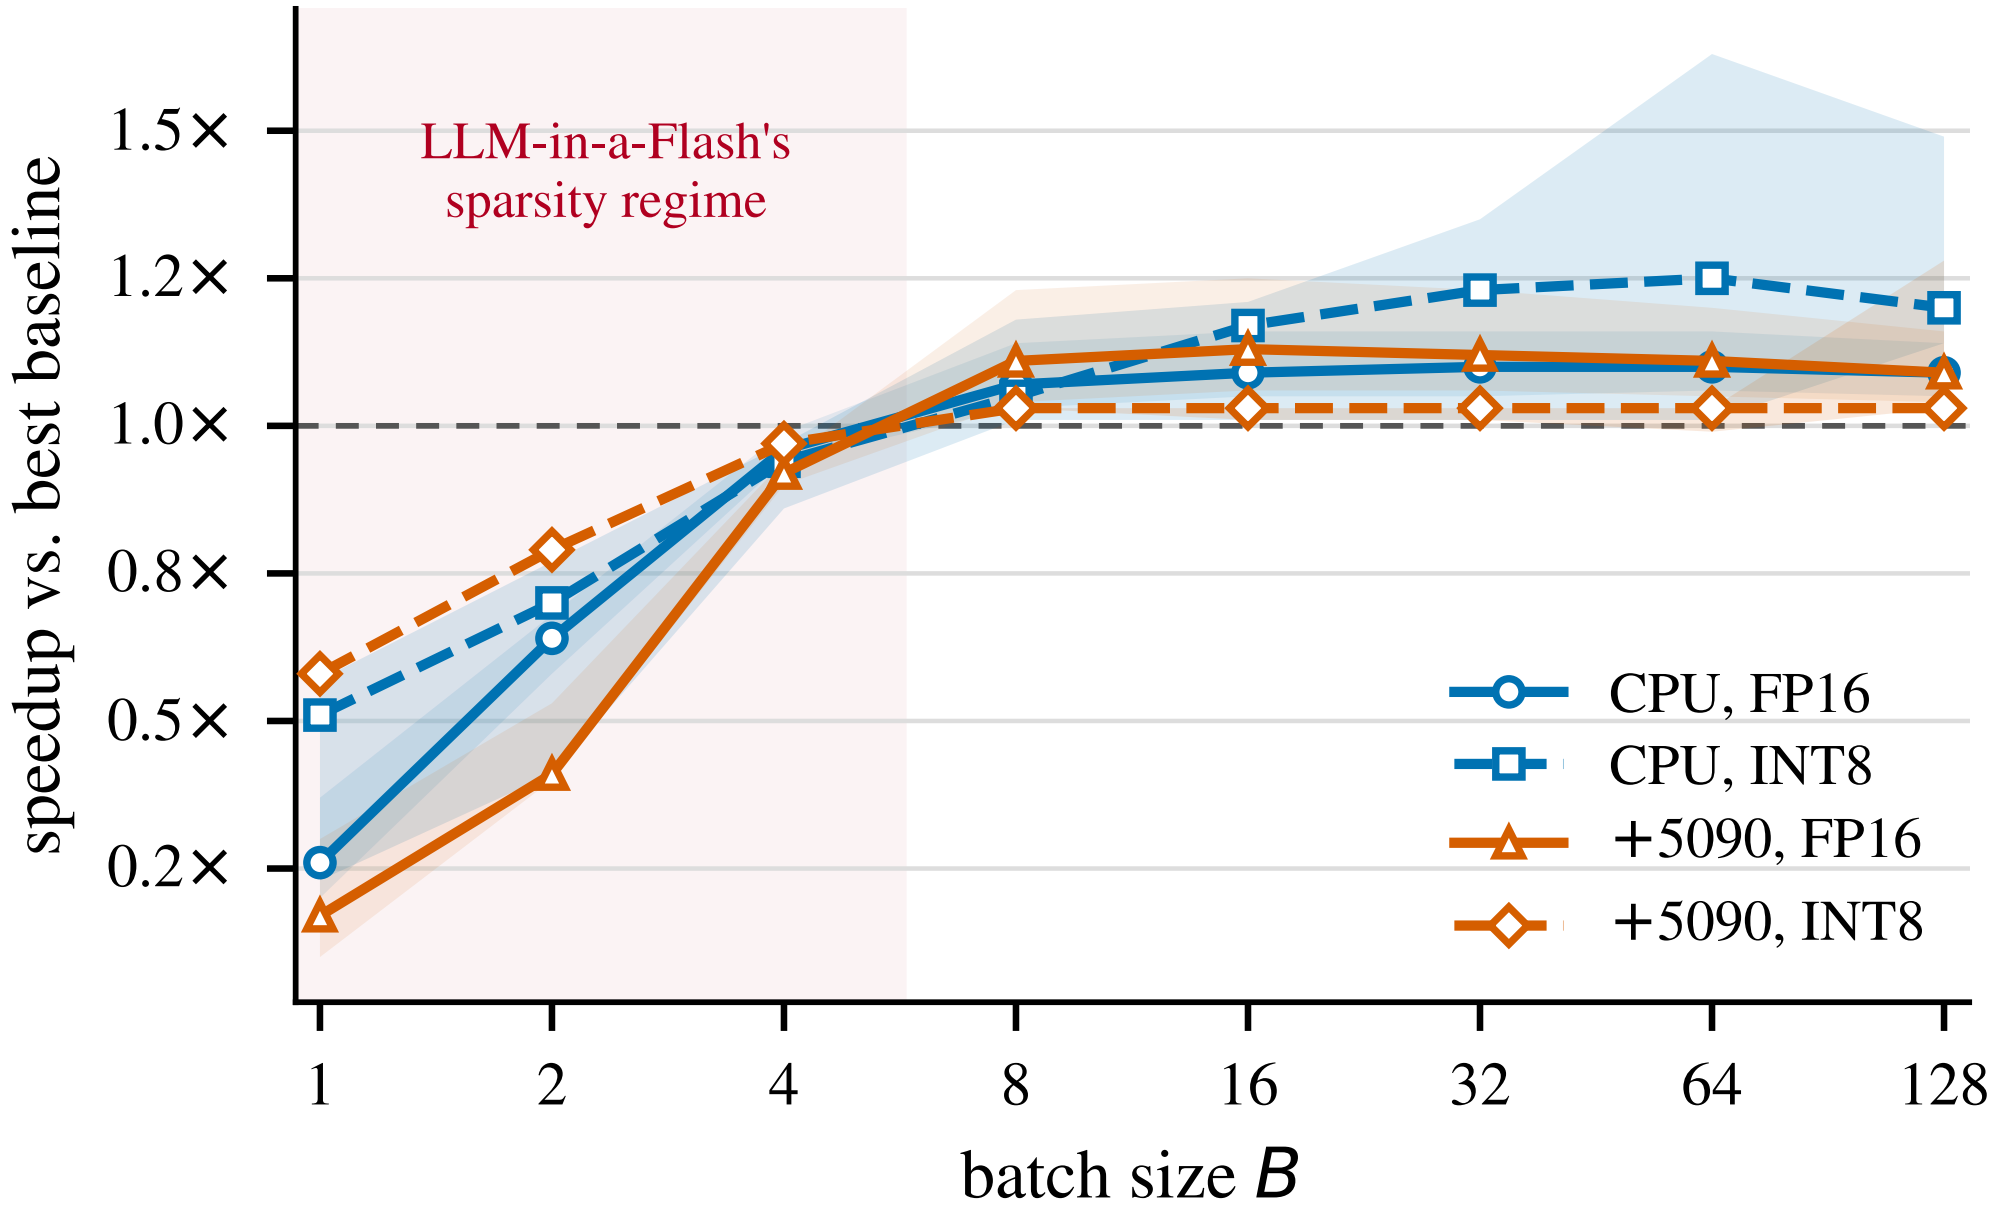
LLM-in-a-Flash's
sparsity regime
speedup vs. best baseline
1.5×
1.2×
1.0×
0.8×
0.5×
0.2×
CPU, FP16
CPU, INT8
+5090, FP16
+5090, INT8
1
2
4
8
16
32
64
128
batch size B