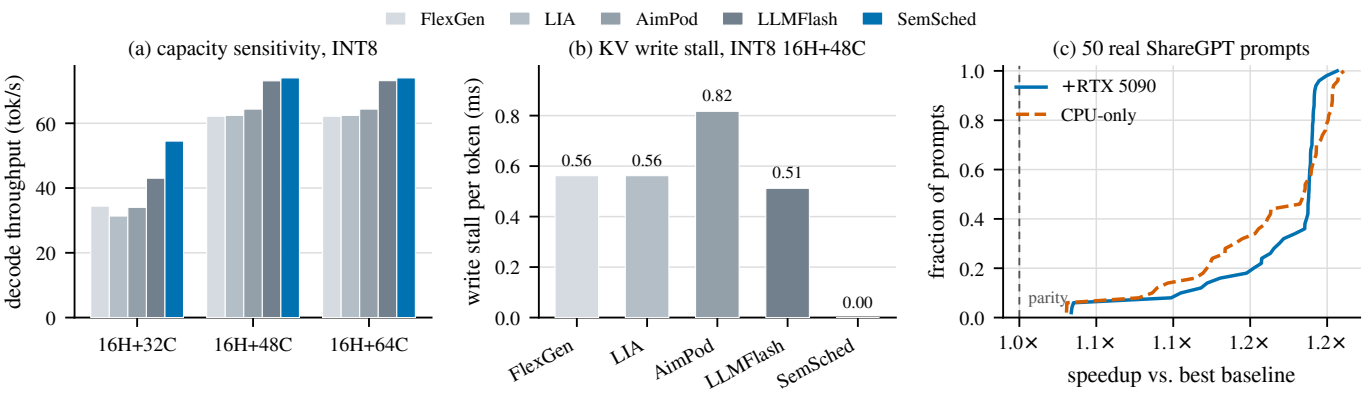
FlexGen
LIA
AimPod
LLMFlash
SemSched
(a) capacity sensitivity, INT8
decode throughput (tok/s)
0
20
40
60
16H+32C
16H+48C
16H+64C
(b) KV write stall, INT8 16H+48C
write stall per token (ms)
0.0
0.2
0.4
0.6
0.8
0.56
0.56
0.82
0.51
0.00
FlexGen
LIA
AimPod
LLMFlash
SemSched
(c) 50 real ShareGPT prompts
fraction of prompts
0.0
0.2
0.4
0.6
0.8
1.0
+RTX 5090
CPU-only
parity
1.0×
1.1×
1.1×
1.2×
1.2×
speedup vs. best baseline