



decode throughput (tok/s)

FlexGen LIA AimPod LLMFlash SemSched

(a) capacity sensitivity, INT8



(b) KV write stall, INT8 16H+48C



(c) 50 real ShareGPT prompts

