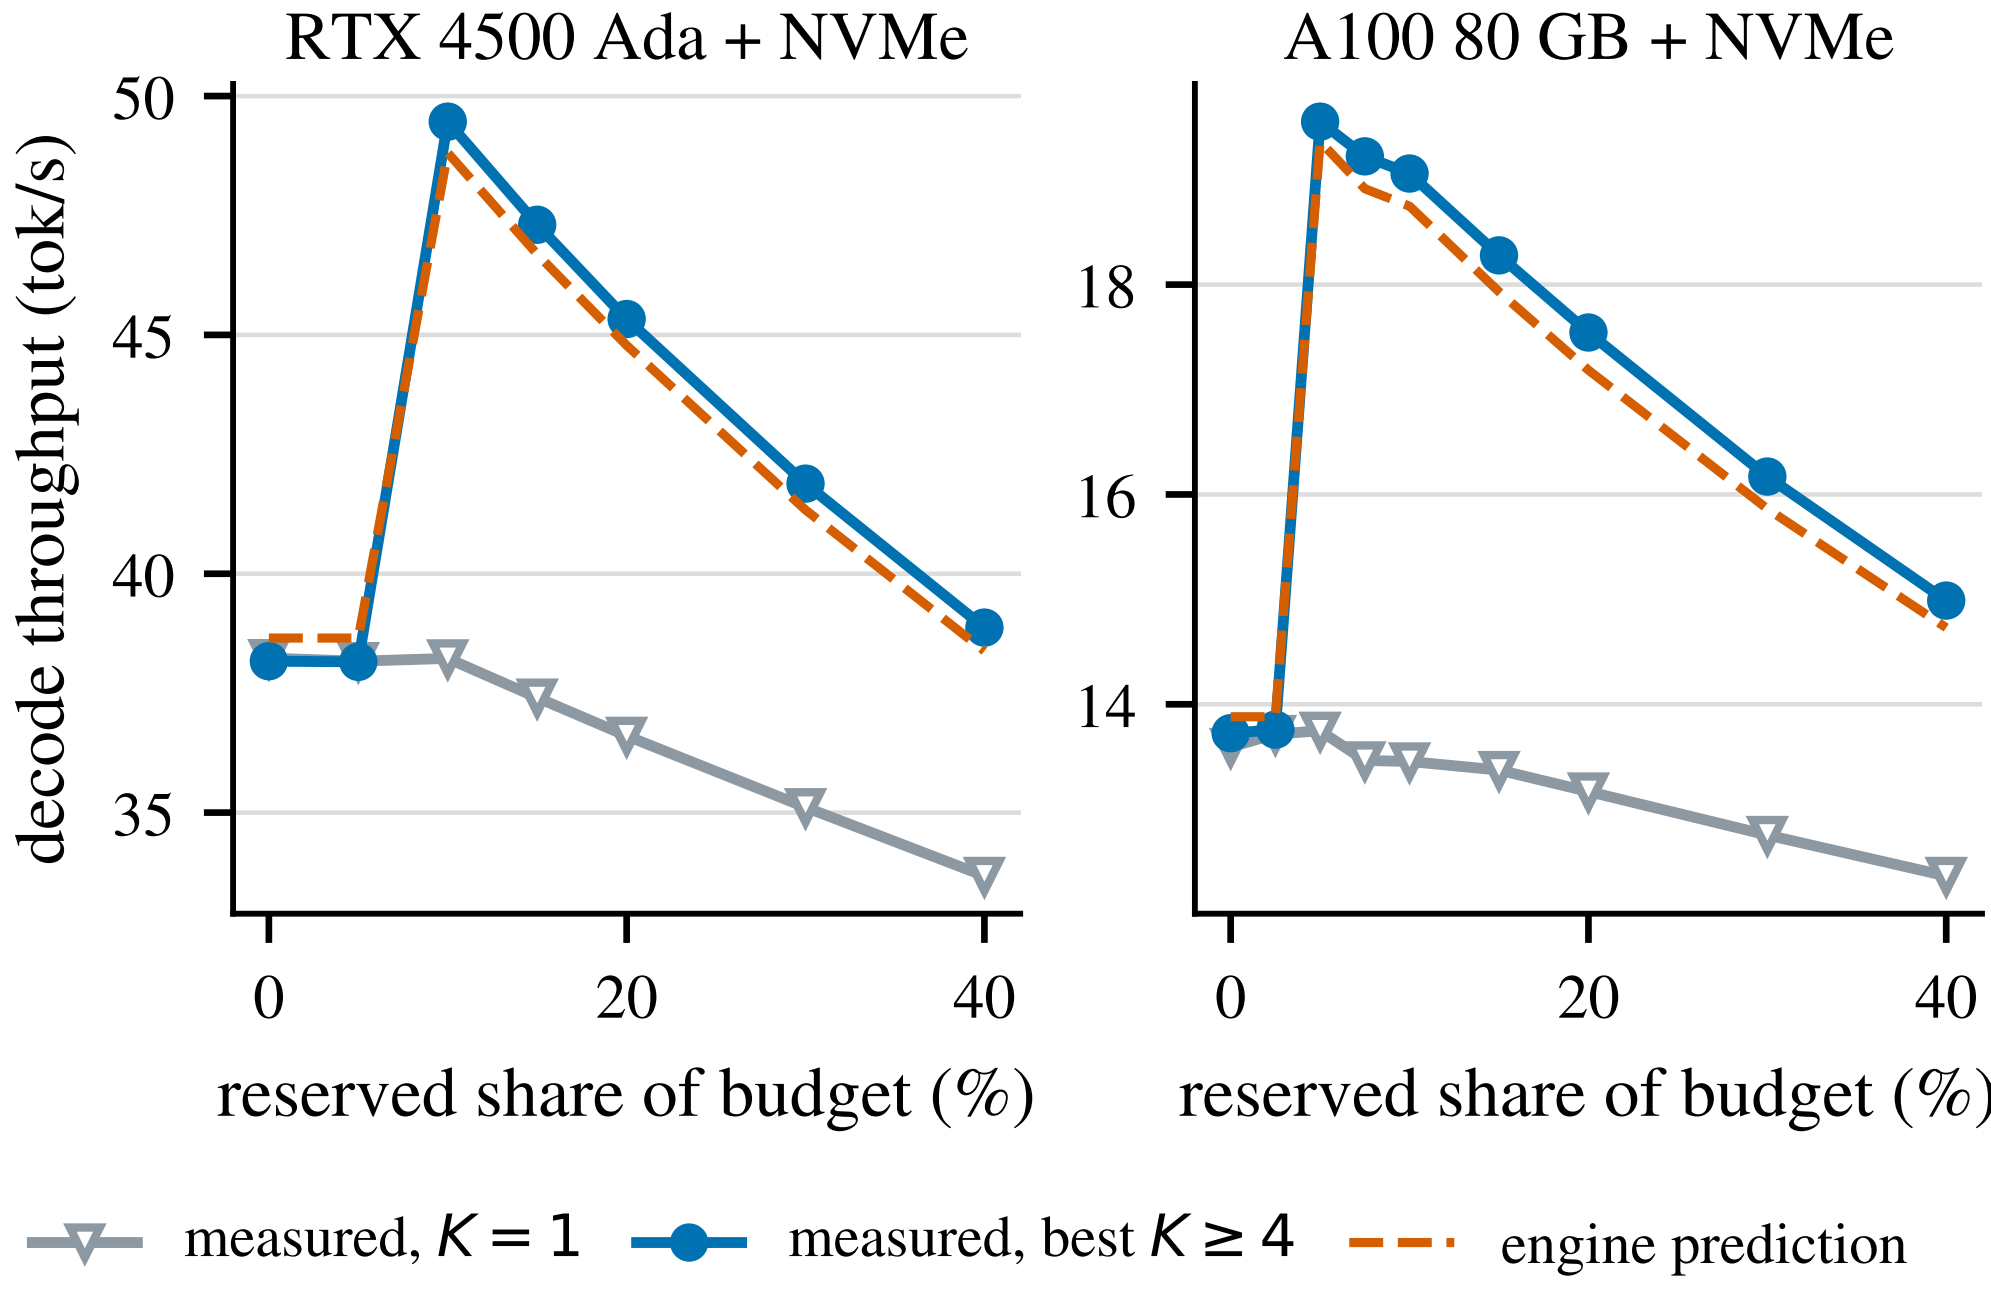
RTX 4500 Ada + NVMe
A100 80 GB + NVMe
decode throughput (tok/s)
50
45
40
35
18
16
14
0
20
40
reserved share of budget (%)
measured, $K = 1$
measured, best $K \geq 4$
engine prediction