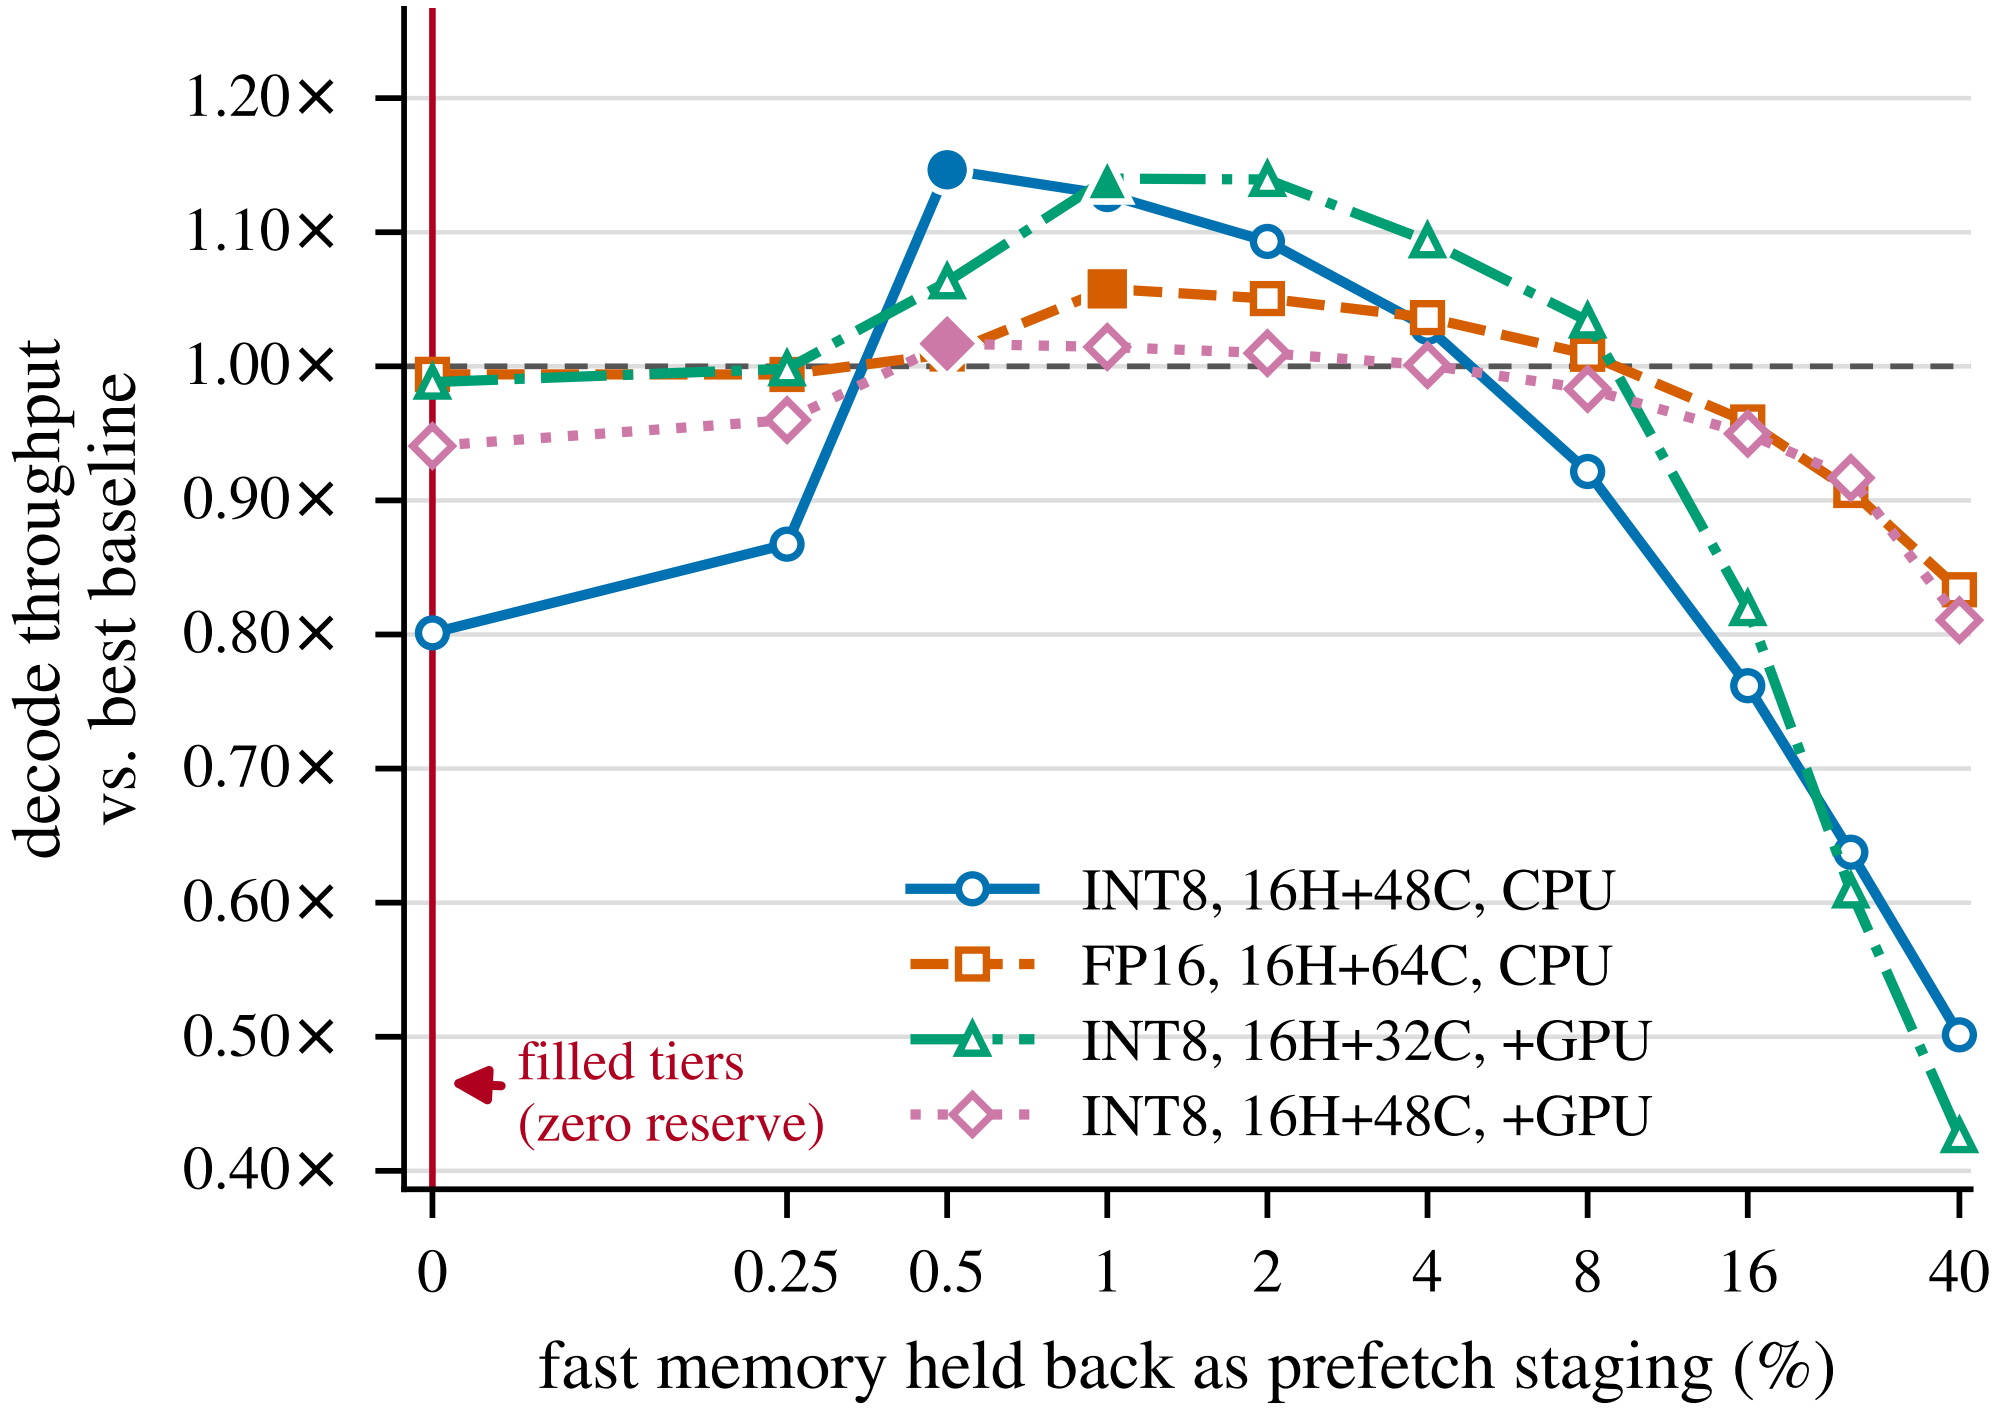
decode throughput
vs. best baseline
1.20×
1.10×
1.00×
0.90×
0.80×
0.70×
0.60×
0.50×
0.40×
filled tiers
(zero reserve)
INT8, 16H+48C, CPU
FP16, 16H+64C, CPU
INT8, 16H+32C, +GPU
INT8, 16H+48C, +GPU
0
0.25
0.5
1
2
4
8
16
40
fast memory held back as prefetch staging (%)